Model accuracy with a common task distribution

Same task weights in every group: the full 2,387-QA task mix. All existing QA retained; no subsampling.

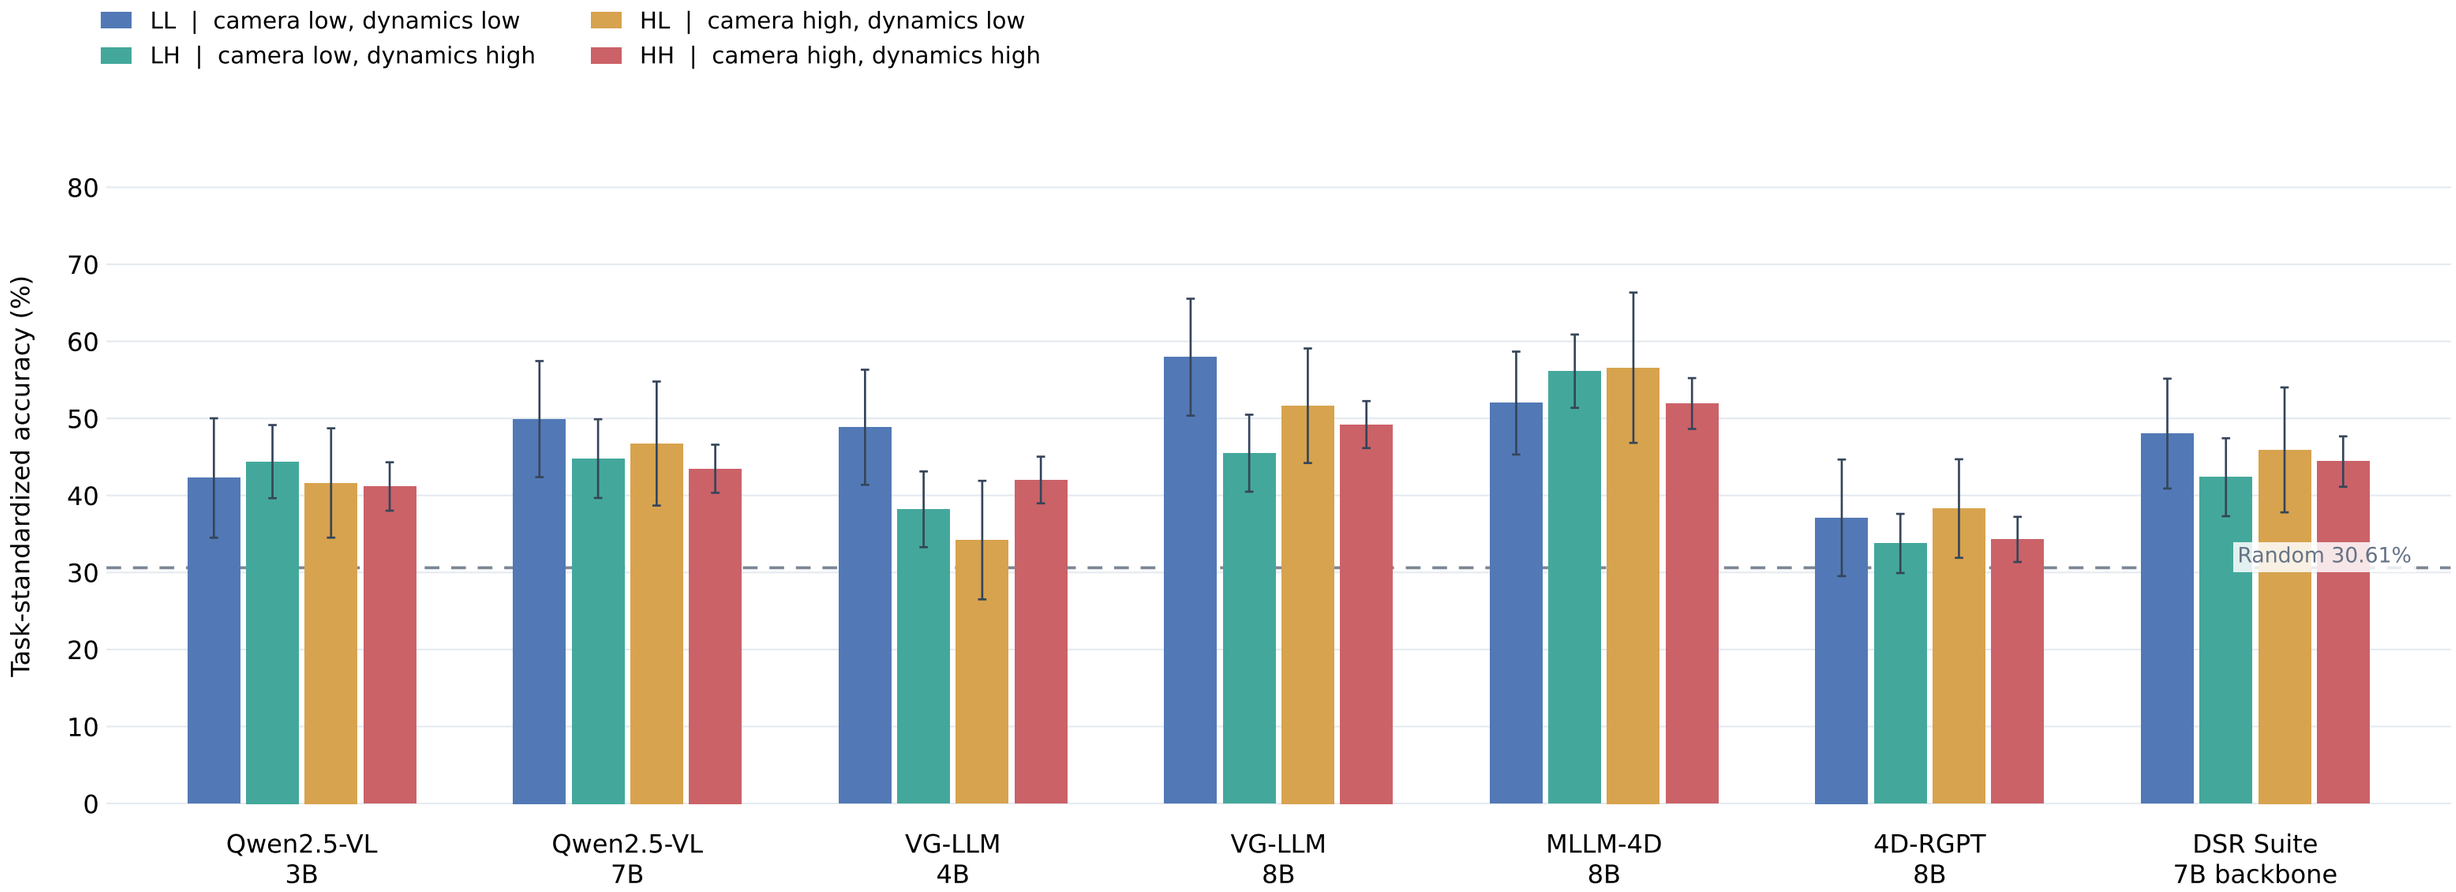

Accuracy (%)

	Qwen2.5-VL 3B	Qwen2.5-VL 7B	VG-LLM 4B	VG-LLM 8B	MLLM-4D 8B	4D-RGPT 8B	DSR Suite 7B backbone
LL	42.26	49.91	48.85	57.94	51.99	37.09	48.03
LH	44.39	44.78	38.20	45.48	56.13	33.76	42.36
HL	41.62	46.73	34.20	51.64	56.58	38.30	45.91
HH	41.16	43.46	42.00	49.20	51.93	34.29	44.40

Task weights T1-T6: 18.77% / 19.35% / 24.88% / 16.97% / 9.80% / 10.22%

Bars: sum of (common task weight x within-group task accuracy). Error bars: video-cluster-robust 95% CIs.

Groups use camera motion x scene dynamics. Camera high = translation high OR rotation high. LL / LH / HL / HH QA: 176 / 848 / 406 / 957.

HL has only 3 T5 and 6 T6 QA; its estimates are less precise. References are automatic geometry annotations. DSR includes spatial modules.