QA-filtered accuracy with a common task distribution

2,262 keep QA / 1,247 videos. All groups use the original 2,387-QA task mix.

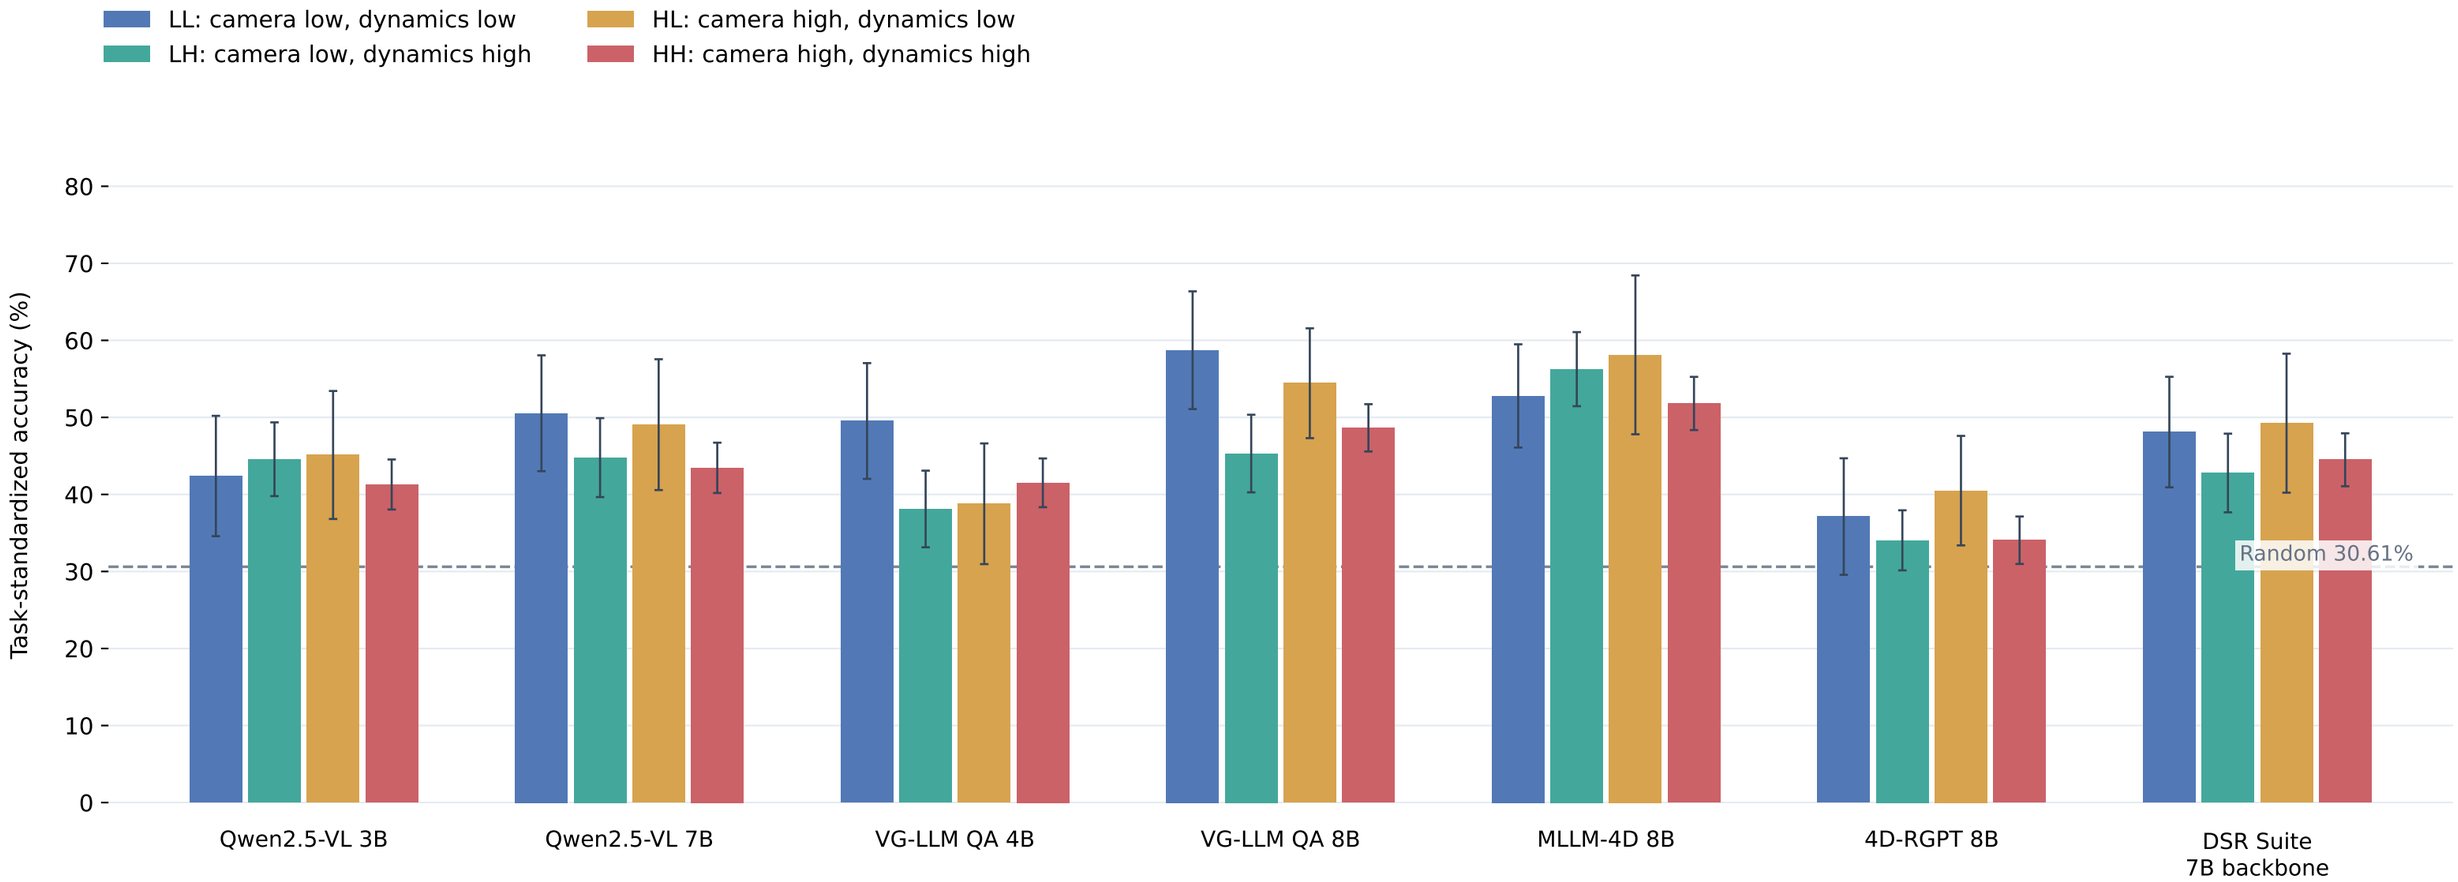

Accuracy (%)

	Qwen2.5-VL 3B	Qwen2.5-VL 7B	VG-LLM QA 4B	VG-LLM QA 8B	MLLM-4D 8B	4D-RGPT 8B	DSR Suite 7B backbone
LL	42.39	50.54	49.53	58.72	52.79	37.13	48.10
LH	44.57	44.78	38.11	45.31	56.26	34.04	42.78
HL	45.12	49.06	38.79	54.43	58.12	40.49	49.25
HH	41.29	43.45	41.51	48.64	51.81	34.05	44.50

T1-T6 common weights: 18.77% / 19.35% / 24.88% / 16.97% / 9.80% / 10.22%

Error bars: video-cluster-robust 95% CIs. Camera high = translation high OR rotation high.

Retained LL / LH / HL / HH QA: 173 / 794 / 399 / 896.

HL has only 3 T5 and 4 T6 QA. Estimates for sparse cells are less precise. References are automatic geometry labels.